

pioSortBed: a multi-mode, parallel BED sorter with native Parquet output

Piotr Balwierz (Genomika.ai)

2026-05-08

Abstract

The BED format is a workhorse of genomic interval analysis, but BED files produced by modern sequencing assays now routinely exceed 10^8 records, making sorting a non-trivial pipeline step. The standard tools — `LC_ALL=C sort -k1,1 -k2,2n`, `bedtools sort`, and `bedops sort-bed` — all rely on single-threaded comparator-based sorting and treat BED records as opaque text. We present `pioSortBed`, a parallel BED sorter built around a chunked memory-mapped parser feeding a least-significant-digit (LSD) radix sort over packed 64-bit (`chromosome-index, position`) keys. Four sort paths share that kernel and let the user trade RAM, wall time, and SSD wear: an in-RAM *classic* path, a two-pass *low-mem-ssd* path, an *external-merge* path with compressed temp runs, and a *multi-pass* path that performs **zero temp-file writes**. On real WGS data (NA12878 GRCh38 100M-read BED, 6.6 GB), `pioSortBed --low-mem-ssd -t 8` sorts in 14.8 s — $12.0\times$ faster than single-threaded GNU `sort`, $3.5\times$ faster than `sort --parallel=8`, $4.4\times$ faster than `bedops sort-bed`, and $9.7\times$ faster than `bedtools sort`. At 500 million synthetic records, `--multi-pass` writes 0 GiB of temp data versus 6.5 GiB for `--external-merge` and 20.2 GiB for either GNU `sort` or `bedops`. An integrated `--bzip --tabix` output mode (htslib-backed) produces a queryable `.bed.gz + .tbi` in a single invocation: on the 100M-record input, `pioSortBed` finishes in 34 s vs 86 s for the canonical `sort | bzip ; tabix` three-tool pipeline, and writes 1.3 GB to disk vs 8.0 GB. `--loci-ss-output` instead emits sorted Parquet conforming to the published LociSSD v2 spec, directly consumable by `polars`, `DuckDB`, `pyranges` / `Loci`, and any Parquet-aware reader, with statistics-based predicate pushdown for fast region queries. At a 10 Mbp region size on the 100M-record BED, `polars + LociSSD` answers a region query in 7.8 ms — $13\times$ faster than `tabix`. The source is a single C++20 file (~3900 lines) released under the MIT license, with a self-contained static Linux x86_64 binary.

Keywords: BED, sorting, genomic intervals, radix sort, parallel, external sort, SSD endurance, Parquet, bioinformatics, file format

Introduction

While the SAM/BAM/CRAM family of binary alignment formats (Danecek et al. 2021) has been rapidly adopted as the standard representation for sequence alignment data, large-scale genomic pipelines continue to rely heavily on BED (Kent et al. 2002) and BED-like

tab-delimited interval formats for downstream data products: peak calls, methylation tracks, candidate regulatory elements, fragment maps, and similar. Reference resources such as the ENCODE project (The ENCODE Project Consortium 2012) distribute much of their derived interval data in BED, and analysis pipelines compose primarily through these files. To support rapid random access on a BED file, three preprocessing steps are typically required: sort, optionally compress (with bgzip (Danecek et al. 2021)), and index (with tabix (Li 2011)). Compression has been heavily parallelised in modern implementations and is no longer a bottleneck; indexing is essentially linear in compressed size and similarly fast. The sort step, performed today by general-purpose tools — GNU `sort -k1,1 -k2,2n` (Free Software Foundation 2025), `bedtools sort` (Quinlan and Hall 2010), and `bedops sort-bed` (Neph et al. 2012) — remains the slowest stage of the three on the inputs characteristic of high-throughput pipelines.

The hardware context has shifted in two ways. First, workhorse pipeline nodes today routinely carry 64–256 GiB of RAM and NVMe-class local SSDs, so the classic external-merge-sort assumption — that the working set must fit in a small RAM budget while spilling intermediate runs to comparatively slow disk — is rarely the binding constraint for ordinary BED inputs. Second, those NVMe SSDs are typically rated for a fixed total bytes written (TBW) endurance — 600 TBW is a common consumer-class number for a 1 TB drive — and pipelines that re-sort the same large BED hundreds of times per year materially erode that budget. A modern BED sorter should be able to exploit on-node parallelism *and* exploit available RAM to avoid temp-file writes when possible, while still degrading gracefully when input genuinely exceeds RAM.

A second, simultaneous shift is in the consumer side of BED data. Pipelines increasingly want sorted BED that is *also* statistically indexed and ready for direct ingest by Parquet-native analytics tools — polars (Vink and Polars contributors 2025), DuckDB (Raasveldt, Mühleisen, and DuckDB Foundation 2025), or pyranges 1 (Stovner and Sætrom 2025). Today this is typically a separate step after sorting (BED-text → BED-to-Parquet), which doubles I/O on multi-gigabyte files. Doing the conversion at sort time, in a single pass, avoids the second read.

All three of the standard sorting tools are effectively single-threaded: `sort --parallel=N` uses multiple threads only inside `std::sort`, not for parsing, and `bedtools sort` and `bedops sort-bed` are serial throughout. All three treat the chromosome name as an opaque string and the start coordinate as a parsed integer compared via a generic comparator. None exploits the fact that BED keys are bounded integers that can be radix-sorted, none parses input in parallel, and none can emit statistically indexed Parquet directly.

Two observations drive the design of `pioSortBed`: (1) the natural BED sort key — the pair (chromosome-index, start-position) — fits in a 64-bit unsigned integer and can be LSD-radix-sorted in $O(n)$ time at very high constant-factor throughput; and (2) BED parsing is embarrassingly parallel over newline-aligned chunks of an `mmap`'d input file. Building on those primitives, `pioSortBed` exposes four sort paths with different RAM / wall time / SSD-wear trade-offs, plus an optional Parquet output mode. We show that the combination yields 5–17 $\times$ speedups over the dominant tools across input sizes from 10^4 to 5×10^8 records, and that the resulting tool is correctness-equivalent to the canonical `LC_ALL=C sort -k1,1 -k2,2n` reference.

Implementation

pioSortBed is a single-source-file C++17 program (~3700 lines, plus the header-only CLI11 (Schreiner and CLI11 contributors 2025) for argument parsing) compiled against the parallel STL backend provided by oneTBB (Intel Corporation 2025) and linked against liblz4 / libzstd (Collet, Kuchera, and Facebook, Inc. 2024). It exposes a single executable that accepts a BED file (uncompressed, gzip-compressed, or read from standard input) and emits sorted BED to standard output, a named file, or a LociSSD Parquet dataset.

Input acquisition

Files on disk are mapped with `mmap` (`PROT_READ | PROT_WRITE | MAP_POPULATE`, with `MADV_HUGEPAGE`), giving the parser zero-copy read-only access to every byte of the input. Gzip-compressed inputs and standard input are slurped into a heap buffer through the same downstream interface, so all subsequent stages operate uniformly on a (base, size) pair. Header lines (`#`, track, browser) are emitted to the output stream in a leading pass and elided from the body.

Parallel parsing

The input is split into newline-aligned chunks, one per worker thread. Each thread parses its chunk into per-chromosome partial linked lists keyed by chromosome name, with read indices that are already global across the file (no rebase pass at merge time). A final serial step concatenates the per-chunk linked lists into a global per-chromosome map.

Chromosome names are stored as (pointer, length) tuples into the input buffer; field 4+ tails are similarly retained as pointers, so per-record metadata is small (24 bytes for the in-RAM paths, 16 bytes for `--low-mem-ssd`) and per-record string copies are avoided in all standard cases. The same parallel parser drives every sort path, including the streaming variants — a property added in v3.5.0 that contributed roughly a 2.4× speedup at 8 threads on `--external-merge` versus the previous serial `getline` parser.

Sorting kernel: radixSort64

The shared sorting kernel is an LSD radix sort over packed 64-bit (`chromosome-index << 32 | position`) keys (Knuth 1998), implemented as 8 passes of 8-bit digits. The implementation skips passes whose corresponding byte is constant across all keys — typical for the high bytes when `chromosome-index` is small and positions fit in 32 bits, so most files require only 4 of the 8 passes. The parallel variant builds per-thread digit histograms in parallel, performs a single serial 2D prefix sum over the (`T threads × 256 buckets`) grid, then scatters in parallel into disjoint output ranges. No atomics are required at any stage. For inputs below ~3 million records, parallel comparator `std::sort` over the same key proves marginally faster and is selected automatically.

Four sort paths

pioSortBed exposes four sort paths, all driven by the same parallel parser and radixSort64 kernel, but differing in how they hold per-record state and how they emit. The defaults are tuned to the common case (a BED file that fits in RAM with room to spare) but the user can override at the command line. Table 1 summarises the trade-offs.

Classic (default). Stores each record in a 24-byte seqread struct (`int beg`, `int end`, plus union slots used during parsing and sorting, plus a `uint16_t lineLen` so output emit can use `fwrite_unlocked` with a known length rather than `strlen`-scanning each line). After parsing, indices into this array are LSD-radix-sorted by the packed 64-bit key, then a linear scan emits each line zero-copy from the `mmap`'d input buffer. Recommended for inputs that comfortably fit in RAM and where wall time matters more than memory footprint.

--low-mem-ssd. A two-pass design. Pass 1 walks the input and builds a flat 16-byte-per-record table (`size_t off`; `uint32_t next`; `uint32_t lineLen`) of per-chromosome linked lists, again in parallel chunks. Pass 2 processes chromosomes concurrently: each chromosome's linked list is materialised into a contiguous index array, sorted with radixSort64 on a single 32-bit position key (no chromosome-index in the key — chromosomes are processed independently), and emitted directly from `mmap` pointers via `fwrite_unlocked`. Output across chromosomes is serialised through a `mutex + condition_variable` barrier so chromosomes appear in the user's chosen order (lexicographic by default, or natural order via `--natural-sort`). The 16-byte index halves the per-record metadata of the classic path. On real-world BED with 24 unevenly sized chromosomes, the per-chromosome parallel emit makes this path the fastest of the four (see §Results).

--external-merge. A bounded-RAM streaming sort for inputs that may exceed available RAM. Pass 1 reads the input in chunks (one chunk per thread, parallel-parsed), fills an in-RAM run buffer up to `--max-mem` bytes, sorts the run with radixSort64, and writes a *compressed* binary run file to a temp directory. Pass 2 opens all run files, maintains a min-heap of (key, runIdx), and emits BED text by popping and refilling the heap. Peak RSS is bounded by `--max-mem` regardless of input size, so the path will sort an arbitrary-sized BED in a fixed memory footprint.

The temp-run codec is selectable via `--merge-codec`; `zstd` is the default. At 200M records / 256 MiB budget, `zstd` is faster *end-to-end* than raw (61.7 s vs 64.0 s), because the disk I/O saved by $\sim 3\times$ compression exceeds the codec's CPU cost. `lz4` trades modest ratio for fast decode. In `WITH_BAM` builds, `htscodecs`' SIMD-vectorised `rANS` order-0 codec is also available and is the fastest in that configuration.

--multi-pass. A K -pass scan that performs **zero temp-file writes**. Pass 1 streams the input via the parallel parser and builds a histogram keyed by (`chromosome-index`, `beg >> 20`) (a 1 MiB position quantum), tracking byte size per bucket. The histogram is bin-packed into K consecutive (`chr`, `beg-quantum`) groups each $\leq \text{--max-mem}$. Passes 2.. $K+1$ re-stream the input, filter to one group, sort + emit. Total writes equal the size of the sorted output (pioSortBed never writes a temp file in this mode); total reads = $(K+1) \times \text{file size}$. The path trades wall time for SSD endurance: it is $2\text{--}3\times$ slower than `--external-merge` at similar RAM budgets but writes nothing the user did not ask for. At input sizes

between $1\times$ and $\sim 5\times$ available RAM — where the alternative is paying 30–100% extra disk writes for external-merge or hitting OOM with the classic path — `--multi-pass` is the right default for repeat-sort workloads on consumer SSDs.

Table 1. Sort-path comparison. *Best for* is the input regime where this path is the right default. *Temp* refers to extra writes beyond the unavoidable sorted output file (i.e. the SSD-wear cost).

Path	Per-record RAM	Temp writes	Best for	Sort orders
classic (default)	24 B + line	none	small–medium, RAM headroom	s, b, 5
--low-mem-ssd	16 B + line	none	$\geq 1\text{M}$ records, SSD-backed	s, b, 5
--external-merge	0 (RAM cap)	$\sim 30\text{--}50\%$ of output (compressed)	input $\gg$ RAM	s only
--multi-pass	0 (RAM cap)	0	$1\times\text{--}5\times$ RAM, wear-sensitive	s only

`--external-merge` and `--multi-pass` currently support only the default coordinate sort (`--sort=s`); `--sort=b` (start+end) and `--sort=5` (strand-aware 5'-end) require the in-RAM paths. Stdin and gzip inputs require classic or low-mem-ssd. Header passthrough, `--natural-sort`, and `--collapse` (where applicable) work uniformly across all four paths.

LociSSD Parquet output

`--lociss-output FILE` makes `pioSortBed` write an Apache Parquet file (Apache Software Foundation 2025b) conforming to the LociSSD v2 spec (Balwierz 2026) instead of (or alongside) a BED text stream. Every file contains the three required loci columns (Chromosome, Start, End) and the derived `MaxEndSoFar` (per-chromosome running maximum of End); standard BED4/5/6/12 inputs additionally emit the typed user columns. The schema is detected from the first record's column count and locked for the whole file:

Input	Additional typed columns (in file order)
BED3	(none)
BED4	Name
BED5	Name, Score
BED6	Strand, Name, Score
BED12	Strand, Name, Score, ThickStart, ThickEnd, ItemRgb, BlockCount, BlockSizes, BlockStarts

Input	Additional typed columns (in file order)
other (7/8/9/10/11/13+)	Tail (catch-all string carrying the raw post-End bytes verbatim)

Column ordering follows FORMAT_SPEC §3.3: required loci columns first, then Strand if present (pulled to position 4), then the remaining user columns in BED order, then MaxEndSoFar last. Score is encoded as string rather than int32 because BED scores can be non-numeric (e.g. .); ThickStart, ThickEnd, BlockCount are int32 per the BED12 specification. The catch-all Tail flavor preserves arbitrary record layouts (narrowPeak’s 10 columns, custom signal-tracking formats with extra fields) and tolerates record-to-record column count variability — useful for messy real-world inputs that the typed flavors would reject. The Parquet footer KV metadata embeds a JSON manifest recording per-chromosome row offsets, value bounds, schema, writer version, and the resolved flavor.

MaxEndSoFar makes the file self-indexing for region queries. A reader issuing (query_chromosome, query_start, query_end) pushes down the predicate

```
Chromosome == query_chromosome
AND Start < query_end
AND End > query_start
AND MaxEndSoFar > query_start
```

to Parquet. The fourth clause is logically redundant with the third but lets a statistics-aware Parquet reader prune row groups via min/max statistics on the *cumulative* MaxEndSoFar column even when a few long intervals would inflate End’s row-group max statistic and defeat pruning of the underlying End column directly. With a well-tuned row_group_size (the spec recommends 65536) this gives row-group precision pruning for free, with no dedicated index scan.

Optional --loci-ss-index (added v3.7.0) embeds a row-precision interval index — an Apache Arrow (Apache Software Foundation 2025a) IPC stream encoding (chromosome, starts, ends, max_end_running, row_id) per chromosome, zstd-compressed and stored under the footer KV key lociSSD_interval_index. The index trades ~24 B/record transient RAM during sort for row-precision pruning at read time. The format spec keeps the index optional; for TB-scale datasets the §7 predicate-pushdown path scales better.

The output integrates directly with downstream Parquet consumers without glue:

```
# polars
import polars as pl
df = pl.read_parquet("regions.loci-ss")
df.filter((pl.col("Chromosome") == "chr1") &
          (pl.col("Start").lt(100_000)))

# DuckDB
import duckdb
```

```
duckdb.sql("SELECT * FROM 'regions.lociss' "
           "WHERE Chromosome = 'chr1' AND Start < 100000")

# Loci (a pyranges 1 derivative with native LociSSD support)
import loci
ds = loci.open_lociss("regions.lociss")
df = ds.region("chr1", 100_000, 200_000)
```

--lociss-output works on every sort path; the corresponding sort-mode restrictions still apply (--sort=b|5 and --collapse are not yet supported with LociSSD output). Cross-path consistency was verified on a 500k-record / 22-chromosome fixture: all four sort paths produce byte-identical Parquet (md5sum matches across runs).

Integrated bgzip + tabix output

--bgzip (and the dependent --tabix) writes the sorted output as a BGZF-compressed .bed.gz (and an adjacent .tbi index) in a single pioSortBed invocation, eliminating the canonical three-tool pipeline `sort | bgzip > out.bed.gz; tabix out.bed.gz` that every production BED workflow ends with. The implementation reuses htslib's BGZF writer (Danecek et al. 2021) driven by a small drainer thread that consumes bytes off a stdout-redirced pipe and feeds them to `bgzf_write`, so every sort path's emit code (which writes BED text to stdout) works unchanged. --tabix then calls `tbx_index_build3` on the finished .bed.gz to write the tabix index in place. The feature is gated on `WITH_HTSLIB=1` at build time and is mutually exclusive with --lociss-output (text vs Parquet output).

Output emission

For BED text output, whichever sort path is chosen, output is reconstructed zero-copy from the original mmap'd input wherever possible. The classic and --low-mem-ssd paths retain a (pointer, length) tuple for each record's line in the input buffer and call `fwrite_unlocked(line, lineLen, 1, stdout)` directly. The classic path writes through an 8 MiB stdout buffer; the --low-mem-ssd path uses a per-chromosome `ChromBuf` that pre-sizes output once and grows arithmetically. The streaming paths reconstruct each output line from the run-file payload (which retains chromosome via a per-run dictionary plus the row's beg, end, strand, and tail bytes).

Results

Hardware and methodology


The headline benchmark hardware is a desktop workstation: 16-core / 32-thread CPU at up to 5 GHz, 62 GiB DDR5, KIOXIA NVMe SSD; Linux 6.x. The synthetic sweep figures (Figures 1, 2) were originally captured on a Lenovo ThinkPad P1 Gen 7 (Intel Core Ultra 7 155H — 16 physical cores / 22 logical threads, 32 GiB DDR5, KIOXIA KXG8AZNV1T02 NVMe SSD) as that rig was where the prior edition of this benchmark was run; cross-rig spot-checks at 100 M records were within 5% of the desktop numbers.

Synthetic fixtures are generated with `tests/gen_fixture.awk` (random chromosome assignments, random `[start, end]` intervals on each chromosome, a `read<i>` name, a numeric score, and a strand flag, in random input order). The real-data benchmark uses NA12878 (HG001) GRCh38 300x Illumina WGS reads from GIAB/NHGRI (Zook et al. 2016), converted with `samtools view | bedtools bamtobed`, then `shuf -n 100000000`. Wall time is measured by `/usr/bin/time -f '%e'`; peak resident set size by `/usr/bin/time -f '%M'`; total bytes written and read by `/usr/bin/time -v`'s File system outputs and inputs fields, which derive from `getrusage(2)`'s `ru_oublock / ru_inblock` counters and capture the total block-layer I/O the process caused, including pages still dirty in the page cache at process exit. Each tool emits its sorted output to a file on the same NVMe device as the input. Correctness is verified per run by `cmp` against the `LC_ALL=C sort -k1,1 -k2,2n` reference; all configurations pass at every size. All benchmarks scripts and CSVs are in `benchmark/` in the repository.


Synthetic sweep, 10^4 to 10^8 records

Figure 1 shows wall time versus input size on a log-x, linear-y scale for the classic and `--low-mem-ssd` paths at 1, 4, and 8 threads, plus GNU sort, `bedops sort-bed`, and `bedtools`. The recommended fast path, `pioSortBed --low-mem-ssd -t 8`, sorts 50 million records in 16.1 s, 100 million in 33.0 s, and 200 million in ~70 s. By comparison, single-threaded `LC_ALL=C sort -k1,1 -k2,2n` requires 1 min 19 s, 2 min 51 s, and 5 min+ at the same sizes — a 4–7× slowdown. Parallel GNU sort (`sort --parallel=8`) recovers most of that gap but still takes 23 s, 48 s, and 117 s — 1.4–1.6× slower than `pioSortBed` at every size in the sweep. `bedops sort-bed` and `bedtools sort`, both effectively single-threaded, range from 2× to 7× slower than the default classic path of `pioSortBed` across the benchmarked range.

Figure 2 shows peak resident set size on the same axes. At 50 million records, `pioSortBed --low-mem-ssd` peaks at 3.0 GiB; the classic path peaks at 4.1 GiB; `bedops sort-bed` (which holds all records in heap-allocated per-record structs and discards the original input bytes after parsing) peaks at 2.7 GiB; `LC_ALL=C sort -k1,1 -k2,2n` at 4.4 GiB single-threaded; `bedtools sort` at 19.4 GiB. At 100 million records `bedtools sort` exhausted memory on a 16 GiB host but completed on the desktop in 49.6 GiB of peak RSS.



Wall time vs. input size, eleven sorter configurations on the same machine. Note the log-x axis: the rightmost point is 10^8 records.



Peak resident set size vs. input size, same configurations as Figure 1. The pioSortBed --low-mem-ssd path scales linearly with input bytes (the mmap'd file plus a 16-byte-per-record index dominate the curve); the classic path adds the per-record seqread array to that floor. bedtools sort's heap-allocated per-record structs scale much worse than any other tool.

Real data: NA12878 100M reads

100 million reads sampled from all standard chromosomes (chr1–22, X, Y) of the GIAB HG001 GRCh38 300x WGS BAM, converted to BED via `bedtools bamtobed`, gives a 6.6 GB BED file with realistic chromosome-size distribution (chr1 is $\sim 5 \times$ chr22). All four


pioSortBed sort paths are benchmarked at -t 1 and -t 8 against the previous pioSortBed version (v3.2.0, two paths only) and the four external tools.

Figure 3 and Table 2 show wall time. The headline result: pioSortBed --low-mem-ssd -t 8 at 14.8 s is the fastest configuration, beating the classic path -t 8 (24.4 s) by 9.6 s. The lead comes from per-chromosome parallel emit: real BED has 24 unevenly sized chromosomes, so the chromosome-parallel low-mem-ssd path keeps every core busy through the emit phase, while the classic path's monolithic sort+emit serialises on the radix sort. The two streaming paths come close — --multi-pass -t 8 at 15.4 s and --external-merge -t 8 at 16.2 s — proving that bounded RAM does not have to mean bounded throughput on workstation hardware.

Table 2. NA12878 GRCh38 100M-read BED (6.6 GB) sort, all configurations. Wall time medians from three runs each.

Tool / mode	Wall time	Peak RSS	Speedup vs GNU sort 1t
pioSortBed --low-mem-ssd -t 8	14.8 s	11.2 GB	11.9×
pioSortBed --multi-pass -t 8	15.4 s	12.0 GB	11.5×
pioSortBed --external-merge -t 8	16.2 s	10.2 GB	10.9×
pioSortBed v3.2.0 classic -t 8	21.2 s	11.8 GB	8.3×
pioSortBed classic -t 8	24.4 s	11.0 GB	7.2×
pioSortBed --sort=b -t 8	25.2 s	10.2 GB	7.0×
pioSortBed --sort=5 -t 8	25.7 s	11.0 GB	6.9×
pioSortBed --low-mem-ssd -t 1	27.1 s	8.2 GB	6.5×
pioSortBed v3.2.0 classic -t 1	28.0 s	11.8 GB	6.3×
pioSortBed classic -t 1	30.2 s	11.0 GB	5.8×
pioSortBed v3.2.0 -low-mem-ssd -t 1	42.7 s	8.4 GB	4.1×
GNU sort --parallel=8	51.8 s	18.5 GB	3.4×
bedops sort-bed	1 min 05.7 s	8.2 GB	2.7×
bedtools sort	2 min 23.8 s	49.6 GB	1.2×

Tool / mode	Wall time	Peak RSS	Speedup vs GNU sort 1t
LC_ALL=C sort - k1,1 -k2,2n	2 min 56.5 s	11.0 GB	1.0×



NA12878 GRCh38 100M-read sort, wall time. Bars colour-coded: pioSortBed v3.7.0 (blue), pioSortBed v3.2.0 baseline (orange), external tools (grey). All -t 8 unless labelled otherwise.

SSD wear: bytes written per sort

Figure 4 plots total bytes written to the block layer (output file + temp files combined) versus input size, for the three streaming-capable configurations of pioSortBed plus bedops and GNU sort, on a synthetic sweep at 16 GiB RAM cap and 8 threads. The output file write (uncompressed sorted BED) is unavoidable and equals the input size. Everything beyond that is temp-file writes — the SSD-wear cost of the sort.


Table 3 quantifies. At 500 million records (21.6 GiB output), --multi-pass writes only the output: **0 GiB of temp data, contributing nothing to NAND wear beyond the result the user actually wanted.** --external-merge with the default zstd codec writes 6.5 GiB of compressed temp runs (~32% of output). bedops sort-bed and GNU sort --parallel=8 both spill to /tmp uncompressed and write 20.2 GiB of temps each — equal to the output, doubling the total disk writes.

Table 3. Temp-file bytes written per sort (output writes excluded), 16 GiB RAM cap, 8 threads.

Reads	extmerge zstd	multipass	bedops	GNU sort -8t
50M	0.65 GiB	0 GiB	0 GiB	0 GiB
100M	1.31 GiB	0 GiB	0 GiB	0 GiB
200M	2.60 GiB	0 GiB	0 GiB	7.99 GiB

Reads	extmerge zstd	multipass	bedops	GNU sort -8t
500M	6.49 GiB	0 GiB	20.23 GiB	20.24 GiB

A practical interpretation: a pipeline that re-sorts a 500M-record BED once a week consumes ~3.4 TBW/year on `--external-merge`, ~10.5 TBW/year on bedops or GNU sort, and 0 TBW/year on `--multi-pass`. On a 1 TB consumer-class SSD with 600 TBW endurance, the difference is non-academic.




Total bytes written to the block layer (output + temp), 16 GiB cap, 8 threads, 50M to 500M records. `--multi-pass` overlaps with bedops and GNU sort lines below 200M because bedops/GNU sort do not yet spill at those sizes; above 200M bedops and GNU sort begin spilling and their total writes diverge upward, while `--multi-pass` continues to track the unavoidable output size.

Streaming-mode scaling

Figure 5 shows wall time at 10M to 500M records on log-log axes for the two pioSortBed streaming paths plus bedops `sort-bed` and GNU `sort --parallel=8`, all at 16 GiB cap and 8 threads. Both pioSortBed streaming paths scale linearly with input size and at 500M reads are 4.7× faster than GNU `sort --parallel=8` and 7.6× faster than bedops `sort-bed`.

The figure also overlays the v3.5.0 baseline (dotted lines) for the two pioSortBed streaming paths. The post-v3.5.0 commit threading `numThreads` through the per-run `radixSort64` call gave a substantial speedup: at 500M records, `--multi-pass` improved from 121.9 s to 80.4 s (−34%), and `--external-merge` from 82.9 s to 66.8 s (−19%).



Streaming-mode scaling, 10M–500M records. v3.7.0 (solid) vs v3.5.0 (dotted) for `--external-merge` and `--multi-pass`; `bedops sort-bed` and `GNU sort --parallel=8` for context.

Query performance and pipeline wall time

The previous subsections measure the *sort* step in isolation. The downstream story is what the sorted BED is actually used for: region queries from analysis pipelines. We compare tabix-indexed bgzipped BED against LociSSD Parquet read by three Parquet-native consumers — polars (Vink and Polars contributors 2025), DuckDB (Raasveldt, Mühleisen, and DuckDB Foundation 2025), and Loci (Stovner and Sætrom 2025) (a pyranges 1 derivative with native LociSSD support, used here for `--lociss-index-aware` queries).

Region-query latency. Figure 6 reports median latency for 333 random queries per region size on the NA12878 100M-record sorted BED, warm cache. tabix wins at small region sizes (median 129 μ s at 1 kbp, 1.2 ms at 100 kbp) thanks to its purpose-built C BED reader and very low per-query setup cost. At the 10 Mbp region size, however, tabix has to scan and parse every BED row in the range and degrades sharply to 105 ms; the Parquet readers’ per-query setup cost ($\sim$ 7–10 ms) becomes amortised across the region, and **polars + LociSSD answers the same query in 7.8 ms — 13 $\times$ faster than tabix**. Loci on LociSSD is also faster than tabix at this size (10.7 ms, 9.8 $\times$). DuckDB is the slowest of the three Parquet readers at 1 kbp and 100 kbp due to per-query SQL parsing overhead and at 10 Mbp returns $\sim$ the same time as tabix.

Table 4. Region-query latency (median, microseconds), 100M-record sorted BED, warm cache, 333 queries per cell.

Region size	tabix bgzip+tabix	polars LociSSD	DuckDB LociSSD	Loci LociSSD
1 kbp	129	7,181	12,372	8,459
100 kbp	1,198	7,079	12,693	8,652
10 Mbp	104,847	7,830	112,994	10,673


The optional `--lociss-index` (row-precision interval index in the Parquet footer) does not help at our test scales: the predicate-pushdown path is already faster than walking the index from Python. At 10M-record files where the index fits comfortably in the footer KV, Loci’s index-on path is $\sim 10\times$ *slower* than its index-off path (25 ms vs 2.5 ms, all sizes). At our 100M-record file the index is multi-GB and exceeds pyarrow’s default thrift footer limit; this is a documented LociSSD-format scaling concern (FORMAT_SPEC.md §6.5.1) that suggests omitting the index at TB scale and relying on predicate pushdown.

Pipeline wall time and bytes written. Table 5 reports the total wall time and total bytes written for four end-to-end recipes that all produce a sorted, queryable BED from the same 100M-record 6.6 GB input. The canonical three-tool pipeline takes 86 s and writes 8 GB to disk. `pioSortBed`’s integrated `--bgzip --tabix` produces the same `.bed.gz + .tbi` in **34 s** while writing only **1.3 GB** (GNU sort spills to temp files, bgzip rewrites everything, tabix walks the bgzipped output). LociSSD goes a step further: `--lociss-output` finishes in **23 s** and writes 1.4 GB — by far the fastest path to a queryable, statistically-indexed file.

Table 5. End-to-end pipeline to produce a sorted, queryable BED from the NA12878 100M-record input. `pioSortBed` runs at `-t 8`.

Recipe	Wall time	Bytes written	Speedup vs canonical
LC_ALL=C sort \ bgzip ; tabix (canonical)	86.1 s	8.0 GB	1.0×
pio --bgzip --tabix (this work)	33.9 s	1.3 GB	2.5×
pio -- lociss- output (this work)	22.9 s	1.4 GB	3.8×
pio -- lociss- output -- lociss-index (this work)	43.9 s	1.9 GB	2.0×

The pipeline-time win combines the sort speedup with eliminating two sequential rewrites (sort $\rightarrow$ bgzip $\rightarrow$ tabix); the bytes-written win combines avoiding GNU sort’s temp-file spills with writing a single compressed output.



Region-query latency on the NA12878 100M-record sorted BED, warm cache, median of 333 queries per cell. tabix queries the bgzip+tabix BED; polars / DuckDB / Loci query the LociSSD Parquet. tabix wins at small region sizes (low per-query setup cost) but degrades to 105 ms at 10 Mbp; polars + LociSSD answers the same 10 Mbp query in 7.8 ms.

Correctness

pioSortBed's BED text output is byte-identical to `LC_ALL=C sort -k1,1 -k2,2n` for default-mode (`--sort=s`) inputs, validated in CI on every commit by a 28-test suite covering BED3, BED6, gzip input, stdin input, the four sort paths, `--collapse`, `--natural-sort`, the `--sort=b` and `--sort=5` orders, and edge cases (empty file, single line, missing trailing newline, header-only file). The benchmark harness re-verifies correctness against the same reference at every size and configuration. For LociSSD output, cross-path consistency was verified on a 500k-record / 22-chromosome fixture: all four sort paths produce byte-identical Parquet (verified by md5sum on the full file payload), and the schema matches the v2 spec exactly.

Discussion

Choosing a sort mode

pioSortBed's four sort paths are not redundant; each is the right default in a different regime. Concretely:

- **Classic** (default) — input fits in RAM with comfortable headroom and wall time is what matters. Smallest constant factor, simplest pipeline story.
- **--low-mem-ssd** — input is $\geq \sim 10^6$ records, fits in RAM, and the node has SSD storage. Half the per-record metadata of classic, and per-chromosome parallel emit makes it the fastest path on real-world multi-chromosome BED.

- **--external-merge** — input genuinely exceeds RAM, or the user wants a fixed RAM budget regardless of input size. Default zstd codec keeps temp-file writes to ~30% of output and is faster end-to-end than uncompressed runs.
- **--multi-pass** — input is in the $1\times\text{--}5\times$ RAM range, the SSD's remaining endurance budget matters, and the user is willing to trade $\sim 2\text{--}3\times$ wall time for zero temp writes. This is the right default for pipelines that re-sort the same BED hundreds of times per year on consumer-class SSDs.

LociSSD as a sort-and-archive output

`--lociss-output` integrates the sort step with the immediate downstream representation that most analytical pipelines actually want. Pipelines that re-read a sorted BED several times pay for re-parsing it from text every time; emitting LociSSD once at sort time replaces those reparses with a Parquet open + statistics-aware predicate evaluation. With typical row-group sizes (65 536 rows) and the `MaxEndSoFar` column, overlap queries prune to the relevant row group(s) without a dedicated index scan. The §Query-performance results back this up: at a 10 Mbp region size on a 100M-record file, polars + LociSSD answers a query in 7.8 ms vs 105 ms for tabix, a $13\times$ win, because tabix must scan every row in the region while Parquet's row-group statistics let polars skip whole groups. The optional `--lociss-index` (row-precision interval index) does *not* help at our test scales — predicate pushdown is already fast enough — but the format keeps it opt-in for the workloads where row-precision pruning is worth its RAM cost. The format itself is a Parquet file conforming to a published, versioned spec (Balwierz 2026); the manifest in the footer KV records writer version and per-chromosome bounds for fast file-level introspection.

Limitations

Coordinates are stored as `int32_t` (≤ 2.15 Gbp per single coordinate) and total record count fits in `uint32_t` ($\leq 4.29 \times 10^9$ records). Both are well above any standard reference genome and any plausible single-file BED, but lifting either limit would require widening the relevant fields throughout the code path.

`--external-merge` and `--multi-pass` currently support only the default coordinate sort (`--sort=s`); `--sort=b` (start+end) and `--sort=5` (strand-aware 5'-end) require the in-RAM paths. `--collapse` (which sums weights for records sharing (chromosome, start)) is currently classic-path-only. Stdin and gzip input are accepted only by the in-RAM paths because the streaming paths require a seekable mmap'd input.

BAM input

An optional in-RAM BAM input mode is available as a build-time flag (`make WITH_BAM=1 HTSLIB=...`); it reuses the same radix sort but reads and writes BAM via htslib (Danecek et al. 2021). On a 1.1 GB / 57.8 M- read BAM, `pioSortBed` produces a record-identical output BAM (`md5sum` match across `samtools view` of `pio.bam`, `samtools sort default.bam`, and `samtools sort -m 4G.bam`). Wall time is competitive but not better

than `samtools sort` (23 s for pio at `-t 8` vs 22 s for `samtools` default vs 18 s for `samtools sort -m 4G`); BGZF compression dominates wall time on BAM and is not the bottleneck `pioSortBed`'s design targets.

`pioSortBed` *does* write less to disk than `samtools sort` at default settings: 1007 MB (output only) vs 2091 MB (output + ~1 GB of temp-file spills). At `samtools sort -m 4G` (large enough to fit the sort in RAM, no spilling) the two tools tie on bytes written but `samtools` uses 3.4× less peak RSS (16 GB vs 24 GB) because it streams records through a partial in-RAM index instead of holding every parsed record. Net: `pioSortBed`'s BAM path is competitive on wall time and bytes written, but is RAM-constrained to BAMs that unpack to roughly half of system RAM (the 1.1 GB compressed BAM above peaked at 24 GB of RSS; multi-GB BAMs would exceed it). We therefore recommend `samtools sort -m N` for large BAM workloads where `samtools`' streaming design fits the available RAM budget, and `pioSortBed` for BED.

Availability

`pioSortBed` is open source under the MIT license. The repository is at <https://github.com/balwierz/pioSortBed>. The current release is v3.8.0. Prebuilt static `pioSortBed-linux-x86_64` binaries are published as GitHub release artefacts; building from source requires `GCC ≥ 10 (C++20)`, `oneTBB`, `lz4`, and `zstd` (`apt install libtbb-dev liblz4-dev libzstd-dev`). Three optional features are opt-in at build time so the default binary has no `htslib` / `Arrow` / `Parquet` dependency: `make WITH_HTSLIB=1` enables `--bgzip` / `--tabix` integrated output (and BAM input under `make WITH_BAM=1 HTSLIB=...`); `make WITH_LOCISS=1` enables `--lociss-output` (requires `Apache Arrow` + `Parquet C++`). The `LociSSD` format specification is in `FORMAT_SPEC.md` in the same repository.

The benchmark scripts, fixtures, gnuplot drivers, and the CSVs underlying every figure and table in this paper are in `benchmark/` in the repository.

Acknowledgements

<TODO: optional — funding, collaborators, beta-testers.>

References

- Apache Software Foundation. 2025a. “Apache Arrow: A Cross-Language Development Platform for in-Memory Analytics.” <https://arrow.apache.org/>.
———. 2025b. “Apache Parquet: A Columnar Storage Format.” <https://parquet.apache.org/>.
Balwierz, Piotr. 2026. “LociSSD File Format Specification, Version 2.” https://github.com/<user>/pioSortBed/blob/master/FORMAT_SPEC.md.

- Collet, Yann, Murray Kucherawy, and Facebook, Inc. 2024. “Zstandard: A Fast Lossless Compression Algorithm.” <https://github.com/facebook/zstd>.
- Danecek, Petr, James K. Bonfield, Jennifer Liddle, John Marshall, Valeriu Ohan, Martin O. Pollard, Andrew Whitwham, et al. 2021. “Twelve Years of SAMtools and BCFtools.” *GigaScience* 10 (2): giab008. <https://doi.org/10.1093/gigascience/giab008>.
- Free Software Foundation. 2025. “GNU Coreutils.” <https://www.gnu.org/software/coreutils/>.
- Intel Corporation. 2025. “oneAPI Threading Building Blocks (oneTBB).” <https://github.com/uxlfoundation/oneTBB>.
- Kent, W. James, Charles W. Sugnet, Terrence S. Furey, Krishna M. Roskin, Tom H. Pringle, Alan M. Zahler, and David Haussler. 2002. “The Human Genome Browser at UCSC.” *Genome Research* 12 (6): 996–1006. <https://doi.org/10.1101/gr.229102>.
- Knuth, Donald E. 1998. *The Art of Computer Programming, Volume 3: Sorting and Searching*. 2nd ed. Addison-Wesley.
- Li, Heng. 2011. “Tabix: Fast Retrieval of Sequence Features from Generic TAB-Delimited Files.” *Bioinformatics* 27 (5): 718–19. <https://doi.org/10.1093/bioinformatics/btq671>.
- Neph, Shane, M. Scott Kuehn, Alex P. Reynolds, Eric Haugen, Robert E. Thurman, Audra K. Johnson, Eric Rynes, et al. 2012. “BEDOPS: High-Performance Genomic Feature Operations.” *Bioinformatics* 28 (14): 1919–20. <https://doi.org/10.1093/bioinformatics/bts277>.
- Quinlan, Aaron R., and Ira M. Hall. 2010. “BEDTools: A Flexible Suite of Utilities for Comparing Genomic Features.” *Bioinformatics* 26 (6): 841–42. <https://doi.org/10.1093/bioinformatics/btq033>.
- Raasveldt, Mark, Hannes Mühleisen, and DuckDB Foundation. 2025. “DuckDB: An Embeddable Analytical Database.” <https://duckdb.org/>.
- Schreiner, Henry F., and CLI11 contributors. 2025. “CLI11: A Header-Only c++11 (and Beyond) Command Line Parser.” <https://github.com/CLIUtils/CLI11>.
- Stovner, Endre Bakken, and Pål Sætrum. 2025. “PyRanges 1: A Clean Modern Python Interface to Genomic Intervals.” *Bioinformatics*.
- The ENCODE Project Consortium. 2012. “An Integrated Encyclopedia of DNA Elements in the Human Genome.” *Nature* 489 (7414): 57–74. <https://doi.org/10.1038/nature11247>.
- Vink, Ritchie, and Polars contributors. 2025. “Polars: A Fast Multi-Threaded, Hybrid-Out-of-Core Query Engine for DataFrames.” <https://pola.rs/>.
- Zook, Justin M., David Catoe, Jennifer McDaniel, Lindsay Vang, Noah Spies, Arend Sidow, Ziming Weng, et al. 2016. “Extensive Sequencing of Seven Human Genomes to Characterize Benchmark Reference Materials.” *Scientific Data* 3: 160025. <https://doi.org/10.1038/sdata.2016.25>.